

Biostat 203B Homework 1

Due Jan 24, 2025 @ 11:59PM

Wenjing Zhou and 806542441

Display machine information for reproducibility:

```
sessionInfo()
```

```
R version 4.4.1 (2024-06-14)
Platform: x86_64-apple-darwin20
Running under: macOS 15.1
```

```
Matrix products: default
```

```
BLAS:   /Library/Frameworks/R.framework/Versions/4.4-x86_64/Resources/lib/libRblas.0.dylib
LAPACK: /Library/Frameworks/R.framework/Versions/4.4-x86_64/Resources/lib/libRlapack.dylib;
```

```
locale:
```

```
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
time zone: America/Los_Angeles
```

```
tzcode source: internal
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

```
loaded via a namespace (and not attached):
```

```
[1] compiler_4.4.1    fastmap_1.2.0     cli_3.6.3         tools_4.4.1
[5] htmltools_0.5.8.1 rstudioapi_0.17.0 yaml_2.3.10       rmarkdown_2.29
[9] knitr_1.48        jsonlite_1.8.9    xfun_0.48         digest_0.6.37
[13] rlang_1.1.4       evaluate_1.0.1
```

Q1. Git/GitHub

No handwritten homework reports are accepted for this course. We work with Git and GitHub. Efficient and abundant use of Git, e.g., frequent and well-documented commits,

is an important criterion for grading your homework.

1. Apply for the [Student Developer Pack](#) at GitHub using your UCLA email. You'll get GitHub Pro account for free (unlimited public and private repositories).
2. Create a **private** repository `biostat-203b-2025-winter` and add Hua-Zhou and TA team (Tomoki-Okuno for Lec 1; `parsajamshidian` and `BowenZhang2001` for Lec 82) as your collaborators with write permission.
3. Top directories of the repository should be `hw1`, `hw2`, ... Maintain two branches `main` and `develop`. The `develop` branch will be your main playground, the place where you develop solution (code) to homework problems and write up report. The `main` branch will be your presentation area. Submit your homework files (Quarto file `qmd`, `html` file converted by Quarto, all code and extra data sets to reproduce results) in the `main` branch.
4. After each homework due date, course reader and instructor will check out your `main` branch for grading. Tag each of your homework submissions with tag names `hw1`, `hw2`, ... Tagging time will be used as your submission time. That means if you tag your `hw1` submission after deadline, penalty points will be deducted for late submission.
5. After this course, you can make this repository public and use it to demonstrate your skill sets on job market.

Solution Done. I've already created a repository named `biostat-203b-2025-winter` and here is the [repository link](#).

Q2. Data ethics training

This exercise (and later in this course) uses the [MIMIC-IV data v3.1](#), a freely accessible critical care database developed by the MIT Lab for Computational Physiology. Follow the instructions at <https://mimic.mit.edu/docs/gettingstarted/> to (1) complete the CITI Data or Specimens Only Research course and (2) obtain the PhysioNet credential for using the MIMIC-IV data. Display the verification links to your completion report and completion certificate here. **You must complete Q2 before working on the remaining questions.** (Hint: The CITI training takes a few hours and the PhysioNet credentialing takes a couple days; do not leave it to the last minute.)

PhysioNet Find Share About News

Access Granted - M

Thank you for signing the data use agreement fo

Solution Done. Here is the screenshot of mimic data access approval.

Q3. Linux Shell Commands

1. Make the MIMIC-IV v3.1 data available at location `~/mimic`. The output of the `ls -l ~/mimic` command should be similar to the below (from my laptop).

```
# content of mimic folder
ls -l ~/mimic/
```

```
total 48
-rw-r--r--@ 1 vickey staff 15199 Oct 10 16:29 CHANGELOG.txt
-rw-r--r--@ 1 vickey staff 2518 Oct 10 17:30 LICENSE.txt
-rw-r--r--@ 1 vickey staff 2884 Oct 11 17:55 SHA256SUMS.txt
drwxr-xr-x@ 23 vickey staff 736 Jan 24 03:06 hosp
drwxr-xr-x@ 11 vickey staff 352 Jan 22 16:08 icu
```

Refer to the documentation <https://physionet.org/content/mimiciv/3.1/> for details of data files. Do **not** put these data files into Git; they are big. Do **not** copy them into your directory. Do **not** decompress the gz data files. These create unnecessary big files and are not big-data-friendly practices. Read from the data folder `~/mimic` directly in following exercises.

Use Bash commands to answer following questions.

2. Display the contents in the folders `hosp` and `icu` using Bash command `ls -l`. Why are these data files distributed as `.csv.gz` files instead of `.csv` (comma separated values) files? Read the page <https://mimic.mit.edu/docs/iv/> to understand what's in each folder.

Solution: Display the contents

```
ls -l ~/mimic/hosp
ls -l ~/mimic/icu
```

```
total 10845000
-rw-r--r--@ 1 vickey staff 19928140 Jun 24 2024 admissions.csv.gz
-rw-r--r--@ 1 vickey staff 427554 Apr 12 2024 d_hcpcs.csv.gz
-rw-r--r--@ 1 vickey staff 876360 Apr 12 2024 d_icd_diagnoses.csv.gz
-rw-r--r--@ 1 vickey staff 589186 Apr 12 2024 d_icd_procedures.csv.gz
-rw-r--r--@ 1 vickey staff 13169 Oct 3 09:07 d_labitems.csv.gz
-rw-r--r--@ 1 vickey staff 33564802 Oct 3 09:07 diagnoses_icd.csv.gz
-rw-r--r--@ 1 vickey staff 9743908 Oct 3 09:07 drgcodes.csv.gz
-rw-r--r--@ 1 vickey staff 811305629 Apr 12 2024 emar.csv.gz
-rw-r--r--@ 1 vickey staff 2162335 Apr 12 2024 hcpcsevents.csv.gz
```

```

-rw-r--r--@ 1 vickey staff 2592909134 Oct 3 09:08 labevents.csv.gz
-rw-r--r--@ 1 vickey staff 117644075 Oct 3 09:08 microbiologyevents.csv.gz
-rw-r--r--@ 1 vickey staff 44069351 Oct 3 09:08 omr.csv.gz
-rw-r--r--@ 1 vickey staff 2835586 Apr 12 2024 patients.csv.gz
-rw-r--r--@ 1 vickey staff 525708076 Apr 12 2024 pharmacy.csv.gz
-rw-r--r--@ 1 vickey staff 666594177 Apr 12 2024 poe.csv.gz
-rw-r--r--@ 1 vickey staff 55267894 Apr 12 2024 poe_detail.csv.gz
-rw-r--r--@ 1 vickey staff 606298611 Apr 12 2024 prescriptions.csv.gz
-rw-r--r--@ 1 vickey staff 7777324 Apr 12 2024 procedures_icd.csv.gz
-rw-r--r--@ 1 vickey staff 127330 Apr 12 2024 provider.csv.gz
-rw-r--r--@ 1 vickey staff 8569241 Apr 12 2024 services.csv.gz
-rw-r--r--@ 1 vickey staff 46185771 Oct 3 09:08 transfers.csv.gz
total 8506784
-rw-r--r--@ 1 vickey staff 41566 Apr 12 2024 caregiver.csv.gz
-rw-r--r--@ 1 vickey staff 3502392765 Apr 12 2024 chartevents.csv.gz
-rw-r--r--@ 1 vickey staff 58741 Apr 12 2024 d_items.csv.gz
-rw-r--r--@ 1 vickey staff 63481196 Apr 12 2024 datatimeevents.csv.gz
-rw-r--r--@ 1 vickey staff 3342355 Oct 3 07:36 icustays.csv.gz
-rw-r--r--@ 1 vickey staff 311642048 Apr 12 2024 ingredientevents.csv.gz
-rw-r--r--@ 1 vickey staff 401088206 Apr 12 2024 inputevents.csv.gz
-rw-r--r--@ 1 vickey staff 49307639 Apr 12 2024 outputevents.csv.gz
-rw-r--r--@ 1 vickey staff 24096834 Apr 12 2024 procedureevents.csv.gz

```

3. Briefly describe what Bash commands `zcat`, `zless`, `zmore`, and `zgrep` do.

Solution

- `zcat` is used to display the contents of a compressed file. It is equivalent to `cat` command for uncompressed files.

```
zcat < ~/mimic/icu/icustays.csv.gz | head
```

```

subject_id,hadm_id,stay_id,first_careunit,last_careunit,intime,outtime,los
10000032,29079034,39553978,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (M
10000690,25860671,37081114,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (M
10000980,26913865,39765666,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (M
10001217,24597018,37067082,Surgical Intensive Care Unit (SICU),Surgical Intensive Care Unit
10001217,27703517,34592300,Surgical Intensive Care Unit (SICU),Surgical Intensive Care Unit
10001725,25563031,31205490,Medical/Surgical Intensive Care Unit (MICU/SICU),Medical/Surgical
10001843,26133978,39698942,Medical/Surgical Intensive Care Unit (MICU/SICU),Medical/Surgical
10001884,26184834,37510196,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (M
10002013,23581541,39060235,Cardiac Vascular Intensive Care Unit (CVICU),Cardiac Vascular Int

```

- `zless` is used to display the contents of a compressed file one screen at a time. It is equivalent to `less` command for uncompressed files.

```
zless ~/mimic/hosp/d_hcpcs.csv.gz | head
```

```
code,category,long_description,short_description
00000,,,Invalid Code
0001F,2,,Composite measures
0002F,2,,Composite measures
0003F,2,,Composite measures
0004F,2,,Composite measures
0005F,2,,Composite measures
0006F,2,,Composite measures
0007F,2,,Composite measures
0008F,2,,Composite measures
```

- `zmore` is used to display the contents of a compressed file one screen at a time. It is equivalent to `more` command for uncompressed files.

```
zmore ~/mimic/hosp/d_hcpcs.csv.gz | head
```

```
code,category,long_description,short_description
00000,,,Invalid Code
0001F,2,,Composite measures
0002F,2,,Composite measures
0003F,2,,Composite measures
0004F,2,,Composite measures
0005F,2,,Composite measures
0006F,2,,Composite measures
0007F,2,,Composite measures
0008F,2,,Composite measures
```

- `zgrep` is used to search the contents of a compressed file using regular expressions. It is equivalent to `grep` command for uncompressed files.

```
zgrep -c "99291" ~/mimic/hosp/services.csv.gz
```

32

4. (Looping in Bash) What's the output of the following bash script?

Solution

```
for datafile in ~/mimic/hosp/{a,l,pa}*.gz
do
    ls -l $datafile
done
```

The output of this bash is displaying all the file names starting with a, l, and pa in the hosp folder.

Display the number of lines in each data file using a similar loop. (Hint: combine linux commands zcat < and wc -l.)

Solution:

```
for datafile in ~/mimic/hosp/{a,l,pa}*.gz
do
    zcat< $datafile | wc -l
done
```

```
546029
158374765
364628
```

5. Display the first few lines of `admissions.csv.gz`. How many rows are in this data file, excluding the header line? Each `hadm_id` identifies a hospitalization. How many hospitalizations are in this data file? How many unique patients (identified by `subject_id`) are in this data file? Do they match the number of patients listed in the `patients.csv.gz` file? (Hint: combine Linux commands zcat <, head/tail, awk, sort, uniq, wc, and so on.)

```
# display the first few lines
zcat < ~/mimic/hosp/admissions.csv.gz | head
```

```
subject_id,hadm_id,admittime,dischtime,deathtime,admission_type,admit_provider_id,admission_
10000032,22595853,2180-05-06 22:23:00,2180-05-07 17:15:00,,URGENT,P49AFC,TRANSFER FROM HOSPI
10000032,22841357,2180-06-26 18:27:00,2180-06-27 18:49:00,,EW EMER.,P784FA,EMERGENCY ROOM,HOS
10000032,25742920,2180-08-05 23:44:00,2180-08-07 17:50:00,,EW EMER.,P19UTS,EMERGENCY ROOM,HOS
10000032,29079034,2180-07-23 12:35:00,2180-07-25 17:55:00,,EW EMER.,P060TX,EMERGENCY ROOM,HOS
10000068,25022803,2160-03-03 23:16:00,2160-03-04 06:26:00,,EU OBSERVATION,P39NWO,EMERGENCY RO
10000084,23052089,2160-11-21 01:56:00,2160-11-25 14:52:00,,EW EMER.,P42H7G,WALK-IN/SELF REFER
10000084,29888819,2160-12-28 05:11:00,2160-12-28 16:07:00,,EU OBSERVATION,P35NE4,PHYSICIAN RI
10000108,27250926,2163-09-27 23:17:00,2163-09-28 09:04:00,,EU OBSERVATION,P40JML,EMERGENCY RO
10000117,22927623,2181-11-15 02:05:00,2181-11-15 14:52:00,,EU OBSERVATION,P47EY8,EMERGENCY RO
```

```
# count the number of rows (excluding the header line)
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | wc -l
```

546028

```
# count the number of unique hospitalizations
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $2}' | sort | uniq | wc -l
```

546028

```
# count the number of unique patients
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $1}' | sort | uniq | wc -l
```

223452

```
# display the first few lines of patients.csv.gz
zcat < ~/mimic/hosp/patients.csv.gz | head
# count the number of unique patients in patients.csv.gz
zcat < ~/mimic/hosp/patients.csv.gz | tail -n +2 | awk -F, '{print $1}' | sort | uniq | wc -l
```

```
subject_id,gender,anchor_age,anchor_year,anchor_year_group,dod
10000032,F,52,2180,2014 - 2016,2180-09-09
10000048,F,23,2126,2008 - 2010,
10000058,F,33,2168,2020 - 2022,
10000068,F,19,2160,2008 - 2010,
10000084,M,72,2160,2017 - 2019,2161-02-13
10000102,F,27,2136,2008 - 2010,
10000108,M,25,2163,2014 - 2016,
10000115,M,24,2154,2017 - 2019,
10000117,F,48,2174,2008 - 2010,
364627
```

Therefore the number of unique patients in `admissions.csv.gz` and `patients.csv.gz` are not the same.

6. What are the possible values taken by each of the variable `admission_type`, `admission_location`, `insurance`, and `ethnicity`? Also report the count for each unique value of these variables in decreasing order. (Hint: combine Linux commands `zcat`, `head/tail`, `awk`, `uniq -c`, `wc`, `sort`, and so on; skip the header line.)

Solution:

```
# display the first few lines of admissions.csv.gz
zcat < ~/mimic/hosp/admissions.csv.gz | head
```

```
subject_id,hadm_id,admittime,dischtime,deathtime,admission_type,admit_provider_id,admission_location
10000032,22595853,2180-05-06 22:23:00,2180-05-07 17:15:00,,URGENT,P49AFC,TRANSFER FROM HOSPITAL
10000032,22841357,2180-06-26 18:27:00,2180-06-27 18:49:00,,EW EMER.,P784FA,EMERGENCY ROOM,HOSPITAL
10000032,25742920,2180-08-05 23:44:00,2180-08-07 17:50:00,,EW EMER.,P19UTS,EMERGENCY ROOM,HOSPITAL
10000032,29079034,2180-07-23 12:35:00,2180-07-25 17:55:00,,EW EMER.,P060TX,EMERGENCY ROOM,HOSPITAL
10000068,25022803,2160-03-03 23:16:00,2160-03-04 06:26:00,,EU OBSERVATION,P39NWO,EMERGENCY ROOM,HOSPITAL
10000084,23052089,2160-11-21 01:56:00,2160-11-25 14:52:00,,EW EMER.,P42H7G,WALK-IN/SELF REFERRAL,HOSPITAL
10000084,29888819,2160-12-28 05:11:00,2160-12-28 16:07:00,,EU OBSERVATION,P35NE4,PHYSICIAN REFERRAL,HOSPITAL
10000108,27250926,2163-09-27 23:17:00,2163-09-28 09:04:00,,EU OBSERVATION,P40JML,EMERGENCY ROOM,HOSPITAL
10000117,22927623,2181-11-15 02:05:00,2181-11-15 14:52:00,,EU OBSERVATION,P47EY8,EMERGENCY ROOM,HOSPITAL
```

```
# count the number of unique values for admission_type
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $6}' | sort | uniq -c |
```

```
177459 EW EMER.
119456 EU OBSERVATION
84437 OBSERVATION ADMIT
54929 URGENT
42898 SURGICAL SAME DAY ADMISSION
24551 DIRECT OBSERVATION
21973 DIRECT EMER.
13130 ELECTIVE
7195 AMBULATORY OBSERVATION
```

```
# count the number of unique values for admission_location
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $8}' | sort | uniq -c |
```

```
244179 EMERGENCY ROOM
163228 PHYSICIAN REFERRAL
56227 TRANSFER FROM HOSPITAL
42365 WALK-IN/SELF REFERRAL
12965 CLINIC REFERRAL
8518 PROCEDURE SITE
6317 TRANSFER FROM SKILLED NURSING FACILITY
5837 INTERNAL TRANSFER TO OR FROM PSYCH
```



```
5734 PACU
402 INFORMATION NOT AVAILABLE
255 AMBULATORY SURGERY TRANSFER
1
```

```
# count the number of unique values for insurance
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $10}' | sort | uniq -c
```

```
244576 Medicare
173399 Private
104229 Medicaid
14006 Other
9355
463 No charge
```

```
# count the number of unique values
zcat < ~/mimic/hosp/admissions.csv.gz | tail -n +2 | awk -F, '{print $13}' | sort | uniq -c
```

```
336538 WHITE
75482 BLACK/AFRICAN AMERICAN
19788 OTHER
13972 WHITE - OTHER EUROPEAN
13870 UNKNOWN
10903 HISPANIC/LATINO - PUERTO RICAN
8287 HISPANIC OR LATINO
7809 ASIAN
7644 ASIAN - CHINESE
6597 WHITE - RUSSIAN
6205 BLACK/CAPE VERDEAN
6070 HISPANIC/LATINO - DOMINICAN
3875 BLACK/CARIBBEAN ISLAND
3495 BLACK/AFRICAN
3478 UNABLE TO OBTAIN
2162 PATIENT DECLINED TO ANSWER
2082 PORTUGUESE
1973 ASIAN - SOUTH EAST ASIAN
1886 WHITE - EASTERN EUROPEAN
1858 HISPANIC/LATINO - GUATEMALAN
1661 ASIAN - ASIAN INDIAN
1526 WHITE - BRAZILIAN
1320 HISPANIC/LATINO - SALVADORAN
```

```

1247 AMERICAN INDIAN/ALASKA NATIVE
920 HISPANIC/LATINO - COLUMBIAN
883 HISPANIC/LATINO - MEXICAN
774 SOUTH AMERICAN
725 HISPANIC/LATINO - HONDURAN
664 ASIAN - KOREAN
641 HISPANIC/LATINO - CUBAN
603 HISPANIC/LATINO - CENTRAL AMERICAN
596 MULTIPLE RACE/ETHNICITY
494 NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER

```

7. The `icustays.csv.gz` file contains all the ICU stays during the study period. How many ICU stays, identified by `stay_id`, are in this data file? How many unique patients, identified by `subject_id`, are in this data file?

Solution:

```

# display the first few lines
zcat < ~/mimic/icu/icustays.csv.gz | head
# the number of ICU stays
zcat < ~/mimic/icu/icustays.csv.gz | tail -n +2 | awk -F, '{print $3}' | sort | uniq | wc -l
# the number of unique patients
zcat < ~/mimic/icu/icustays.csv.gz | tail -n +2 | awk -F, '{print $1}' | sort | uniq | wc -l

```

```

subject_id,hadm_id,stay_id,first_careunit,last_careunit,intime,outtime,los
10000032,29079034,39553978,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (MICU)
10000690,25860671,37081114,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (MICU)
10000980,26913865,39765666,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (MICU)
10001217,24597018,37067082,Surgical Intensive Care Unit (SICU),Surgical Intensive Care Unit (SICU)
10001217,27703517,34592300,Surgical Intensive Care Unit (SICU),Surgical Intensive Care Unit (SICU)
10001725,25563031,31205490,Medical/Surgical Intensive Care Unit (MICU/SICU),Medical/Surgical Intensive Care Unit (MICU/SICU)
10001843,26133978,39698942,Medical/Surgical Intensive Care Unit (MICU/SICU),Medical/Surgical Intensive Care Unit (MICU/SICU)
10001884,26184834,37510196,Medical Intensive Care Unit (MICU),Medical Intensive Care Unit (MICU)
10002013,23581541,39060235,Cardiac Vascular Intensive Care Unit (CVICU),Cardiac Vascular Intensive Care Unit (CVICU)
94458
65366

```

The number of ICU stays is 94458 and the number of unique patients is 65366 in this data file.

8. *To compress, or not to compress. That's the question.* Let's focus on the big data file `labevents.csv.gz`. Compare compressed gz file size to the uncompressed file size. Compare the run times of `zcat < ~/mimic/labevents.csv.gz | wc -l` versus

`wc -l labevents.csv`. Discuss the trade off between storage and speed for big data files. (Hint: `gzip -dk < FILENAME.gz > ./FILENAME`. Remember to delete the large `labevents.csv` file after the exercise.)

Solution:

```
# compare compressed and uncompressed file size
ls -lh ~/mimic/hosp/labevents.csv.gz
# Uncompress the file temporarily and check its size
gzip -dk < ~/mimic/hosp/labevents.csv.gz > ./labevents.csv
ls -lh labevents.csv
```

```
-rw-r--r--@ 1 vickey  staff   2.4G Oct  3 09:08 /Users/vickey/mimic/hosp/labevents.csv.gz
-rw-r--r--@ 1 vickey  staff    17G Jan 24 14:23 labevents.csv
```

```
# compare run times
time zcat < ~/mimic/hosp/labevents.csv.gz | wc -l
time wc -l ./labevents.csv
```

158374765

```
real    0m35.699s
user    0m51.265s
sys 0m3.174s
158374765 ./labevents.csv
```

```
real    0m39.578s
user    0m34.726s
sys 0m4.338s
```

The compressed file size is smaller than the uncompressed file size. The run time of `zcat < ~/mimic/labevents.csv.gz | wc -l` is faster than `wc -l ~/mimic/labevents.csv`. The trade off between storage and speed for big data files is that compressed files save storage space but take longer time to read and write, while uncompressed files take more storage space but are faster to read and write.

Q4. Who's popular in *Pride and Prejudice*

1. You and your friend just have finished reading *Pride and Prejudice* by Jane Austen. Among the four main characters in the book, Elizabeth, Jane, Lydia, and Darcy, your friend thinks that Darcy was the most mentioned. You, however, are certain it was

Elizabeth. Obtain the full text of the novel from <http://www.gutenberg.org/cache/epub/42671/pg42671.txt> and save to your local folder.

```
wget -nc http://www.gutenberg.org/cache/epub/42671/pg42671.txt
```

Explain what `wget -nc` does. Do **not** put this text file `pg42671.txt` in Git. Complete the following loop to tabulate the number of times each of the four characters is mentioned using Linux commands.

Solution: `wget -nc` is used to download the file only if it does not exist in the current directory.

```
wget -nc http://www.gutenberg.org/cache/epub/42671/pg42671.txt

for char in Elizabeth Jane Lydia Darcy
do
    echo $char:
    # some bash commands here
    grep -oi "$char" pg42671.txt | wc -l
done
```

2. What's the difference between the following two commands?

```
echo 'hello, world' > test1.txt
```

and

```
echo 'hello, world' >> test2.txt
```

Solution: The difference of the two commands is that the first command `echo 'hello, world' > test1.txt` writes the output to a file `test1.txt` and overwrites the file if it already exists, while the second command `echo 'hello, world' >> test2.txt` appends the output to a file `test2.txt` if the file already exists. `test2.txt` will have the content `hello, world` twice if I run the second command twice.

3. Using your favorite text editor (e.g., `vi`), type the following and save the file as `middle.sh`:

```
#!/bin/sh
# Select lines from the middle of a file.
# Usage: bash middle.sh filename end_line num_lines
head -n "$2" "$1" | tail -n "$3"
```

Using `chmod` to make the file executable by the owner, and run

```
chmod +x middle.sh
```

```
./middle.sh pg42671.txt 20 5
```

Explain the output. Explain the meaning of "\$1", "\$2", and "\$3" in this shell script. Why do we need the first line of the shell script?

Solution: The meaning of "\$1", "\$2", and "\$3" in this shell script is the first, second, and third argument passed to the script, respectively. The first line of the shell script `#!/bin/sh` is called a shebang line, which tells the system which interpreter to use to execute the script. In this case, it tells the system to use the `sh` interpreter to execute the script.

Q5. More fun with Linux

Try following commands in Bash and interpret the results: `cal`, `cal 2025`, `cal 9 1752` (anything unusual?), `date`, `hostname`, `arch`, `uname -a`, `uptime`, `who am i`, `who, w`, `id`, `last | head`, `echo {con,pre}{sent,fer}{s,ed}`, `time sleep 5`, `history | tail`.

Solution

```
cal
```

```
      January 2025
Su Mo Tu We Th Fr Sa
                1  2  3  4
 5  6  7  8  9 10 11
12 13 14 15 16 17 18
19 20 21 22 23 24 25
26 27 28 29 30 31
```

```
cal 2025
```

```

                2025
      January      February      March
Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa
                1  2  3  4                1                1
 5  6  7  8  9 10 11 2  3  4  5  6  7  8 2  3  4  5  6  7  8
12 13 14 15 16 17 18 9 10 11 12 13 14 15 9 10 11 12 13 14 15
```

19	20	21	22	23	24	25	16	17	18	19	20	21	22	16	17	18	19	20	21	22
26	27	28	29	30	31	23	24	25	26	27	28	23	24	25	26	27	28	29		
												30	31							

April							May							June						
Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa
		1	2	3	4	5					1	2	3	1	2	3	4	5	6	7
6	7	8	9	10	11	12	4	5	6	7	8	9	10	8	9	10	11	12	13	14
13	14	15	16	17	18	19	11	12	13	14	15	16	17	15	16	17	18	19	20	21
20	21	22	23	24	25	26	18	19	20	21	22	23	24	22	23	24	25	26	27	28
27	28	29	30				25	26	27	28	29	30	31	29	30					

July							August							September						
Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa
		1	2	3	4	5						1	2		1	2	3	4	5	6
6	7	8	9	10	11	12	3	4	5	6	7	8	9	7	8	9	10	11	12	13
13	14	15	16	17	18	19	10	11	12	13	14	15	16	14	15	16	17	18	19	20
20	21	22	23	24	25	26	17	18	19	20	21	22	23	21	22	23	24	25	26	27
27	28	29	30	31			24	25	26	27	28	29	30	28	29	30				
							31													

October							November							December						
Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa	Su	Mo	Tu	We	Th	Fr	Sa
			1	2	3	4							1		1	2	3	4	5	6
5	6	7	8	9	10	11	2	3	4	5	6	7	8	7	8	9	10	11	12	13
12	13	14	15	16	17	18	9	10	11	12	13	14	15	14	15	16	17	18	19	20
19	20	21	22	23	24	25	16	17	18	19	20	21	22	21	22	23	24	25	26	27
26	27	28	29	30	31		23	24	25	26	27	28	29	28	29	30	31			
							30													

cal 9 1752

September 1752

Su	Mo	Tu	We	Th	Fr	Sa
		1	2	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30

Unusual thing: The days from 3rd to 13th on Sepetember of 1752 in the calnedar are missing.

```
date
```

```
Fri Jan 24 14:24:52 PST 2025
```

```
hostname
```

```
MacBookPro.lan
```

```
arch
```

```
i386
```

```
uname -a
```

```
Darwin MacBookPro.lan 24.1.0 Darwin Kernel Version 24.1.0: Thu Oct 10 21:02:27 PDT 2024; root
```

```
uptime
```

```
14:24 up 11:21, 1 user, load averages: 2.98 2.64 2.50
```

```
who am i
```

```
vickey Jan 24 14:24
```

```
who
```

```
vickey console Jan 24 03:04
```

```
w
```

```
14:24 up 11:21, 1 user, load averages: 2.98 2.64 2.50
```

USER	TTY	FROM	LOGIN@	IDLE	WHAT
vickey	console	-	3:04	11:20	-

```
id
```

```
uid=501(vickey) gid=20(staff) groups=20(staff),101(access_bpf),12(everyone),61(localaccounts)
```

```
last | head
```

```
wtmp begins Fri Jan 24 14:24:52 PST 2025
```

```
echo {con,pre}{sent,fer}{s,ed}
```

```
consents consented confers conferred presents presented prefers preferred
```

```
time sleep 5
```

```
real    0m5.011s
user    0m0.002s
sys     0m0.006s
```

```
history | tail
```

```
9 brew install wget 10 history | tail 11 git branch 12 git check out main 13 git checkout main
14 ls 15 git branch 16 git checkout main 17 git merge develop 18 history | tail
```

Q6. Book

1. Git clone the repository <https://github.com/christophergandrud/Rep-Res-Book> for the book *Reproducible Research with R and RStudio* to your local machine. Do **not** put this repository within your homework repository `biostat-203b-2025-winter`.
2. Open the project by clicking `rep-res-3rd-edition.Rproj` and compile the book by clicking `Build Book` in the `Build` panel of RStudio. (Hint: I was able to build `git_book` and `epub_book` directly. For `pdf_book`, I needed to add a line `\usepackage{hyperref}` to the file `Rep-Res-Book/rep-res-3rd-edition/latex/preabmle.tex`.)

The point of this exercise is (1) to obtain the book for free and (2) to see an example how a complicated project such as a book can be organized in a reproducible way. Use `sudo apt install PKGNAME` to install required Ubuntu packages and `tlmgr install PKGNAME` to install missing TeXLive packages.

For grading purpose, include a screenshot of Section 4.1.5 of the book here.

4.1.5 Spaces in directory and file

It is good practice to avoid putting spaces in file and directory names. For example, I called the example in Figure 4.1 “example-project” rather than “Example Project” and directory names can sometimes create problems for programs trying to read the file path. The presence of a space in the path name indicates that the path name has spaces and is not easily readable without using the quote convention.

One approach is to use a convention that makes clear the naming convention you are using. A common convention is to make it clear if something is an R object or a function. If we adopt the underscore method for R objects, then for example (e.g. `health_data`) we could use hyphenated names for functions. For example: `example-source`, `kebab-case`.

Solution: Here is the screenshot of Section 4.1.5 of the book.